

Random Forest

Random Forest is a popular machine learning algorithm used for both classification and regression tasks. It's an ensemble learning method, which means it combines multiple simpler models to create a more powerful and robust predictive model.

Ensemble of Decision Trees

- A Random Forest is made up of many individual decision trees.
- Each tree is trained on a random subset of the data and features.

Bagging (Bootstrap Aggregating)

- Random Forest uses a technique called bagging to create diverse trees.
- For each tree, a random sample of the training data is selected with replacement (bootstrap sampling).

Feature Randomness

- At each node split in a tree, only a random subset of features is considered.
- This introduces additional randomness and helps to decorrelate the trees.

Voting/Averaging

- For classification tasks, each tree "votes" for a class, and the majority vote wins.
- For regression tasks, the predictions of all trees are averaged.

Hyperparameters

- Number of trees in the forest.
- Maximum depth of trees.
- Minimum number of samples required to split an internal node.
- Number of features to consider for the best split.

Out-of-Bag (OOB) Error

- Random Forest can use the samples not included in the bootstrap sample for each tree to estimate the model's performance without a separate validation set.

Advantages

- Reduces overfitting compared to individual decision trees.
- Handles high-dimensional data well.
- Can capture complex interactions between features.
- Provides feature importance rankings.
- Robust to outliers and noise.

