

# KNN

K-Nearest Neighbors (KNN) is a simple, versatile, and widely used machine learning algorithm for both classification and regression tasks.

## Distance Metrics

Common distance measures include:

- Euclidean distance (most common)
- Manhattan distance
- Minkowski distance
- Hamming distance (for categorical variables)

### 1. Distance Calculation:


The core of KNN is calculating distances between data points. The most common distance metric is Euclidean distance, but others can be used.

Euclidean distance between two points  $p$  and  $q$  in  $n$ -dimensional space:

$$d(p, q) = \sqrt{[(p_1 - q_1)^2 + (p_2 - q_2)^2 + \dots + (p_n - q_n)^2]}$$

Or more concisely:

$$d(p, q) = \sqrt{[\sum_i (p_i - q_i)^2]}$$



## K-Nearest Neighbor (KNN)

## KNN for Regression:

Process: a. For a new data point, find the K nearest neighbors in the training set. b. Calculate the average (or weighted average) of the target values of these K neighbors. c. Use this average as the prediction for the new data point.

## Steps in regression

Let  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$  be the training set, where  $x$  are feature vectors and  $y$  are target values.

For a new point  $x$ :

- Calculate distances to all training points:  $d(x, x_i)$  for  $i = 1$  to  $n$
- Select K nearest neighbors. Let  $S$  be the set of these K neighbors.
- Prediction is the average of the K neighbors' target values:

$$y_{pred} = (1/K) * \sum(y_i)$$

in S Weighted version:

$$y_{pred} = \sum(w_i * y_i) / \sum(w_i)$$

where  $w_i = 1 / d(x, x_i)^2$

## KNN for Classification:

The process is similar, but instead of averaging, we use majority voting.

Let  $C$  be the set of classes, and  $c(x_i)$  be the class of point  $x_i$ .

## Steps in classification

- Calculate distances as in regression.
- Select K nearest neighbors (set  $S$ ).
- For each class  $j$  in  $C$ , count occurrences in  $S$ :  $\text{count}_j = \sum I(c(x_i) = j)$ , for all  $x_i$  in  $S$  where  $I$  is the indicator function (1 if true, 0 if false)

d. Predict the class with the highest count:  $y_{\text{pred}} = \text{argmax}_j(\text{count}_j)$

## Advantages:

- Simple to understand and implement
- No assumptions about data distribution (non-parametric)
- Can be used for both classification and regression
- Works well with multi-class problems

## Disadvantages:

- Computationally expensive for large datasets
- Sensitive to irrelevant features and the scale of the data
- Requires feature scaling
- Does not work well with high-dimensional data (curse of dimensionality)